

FactoryLM / MIRA - Strategy Brief

The wedge, the competitive map, and the path to Provelt! 2027 - June 22, 2026

1. The wedge - the one thing

FactoryLM is the maintenance-context layer that makes a factory's messy reality - PLC tags, manuals, CMMS, tribal knowledge, live telemetry - trustworthy enough for AI, one asset at a time, human-approved, on top of ANY Unified Namespace. MIRA is the grounded agent that proves the context is real by diagnosing with cited sources.

The product is the context. The agent is the proof. Lead with the context platform, never with the copilot: 'AI maintenance copilot' is a crowded category we lose on distribution. 'The layer that makes any maintenance AI trustworthy on YOUR messy factory' is the empty chair.

The adapters stay. Slack, Telegram, the Ignition 'Ask MIRA' panel, QR-scan, web - all valued consumption surfaces. The only rule: every adapter renders the same approved-context answer (same citations, same score, same read-only guarantee). 'Lead with context, not copilot' is about the pitch, not the plumbing.

2. Why this wins - the competitive map

The market splits into two camps, and neither owns the chair between them.

Platforms / plumbing (Fuuz, HighByte, Ignition, Cognite data layer)

Own: UNS transport, ISA-95 hierarchy, MES / WMS / OEE, connectors. Assume away: the maintenance BRAIN. Fuuz - the one you flagged - is ~43 people, bootstrapped, mid-market, and DELIBERATELY builds no diagnostic AI ('bring your own LLM via GraphQL'). A partner/integration target, not a rival.

Copilots (MaintainX, Cognite, AVEVA, C3, UptimeAI)

Own: grounded RAG over manuals and work orders. Assume away: the CONTEXTUALIZATION - they answer AFTER the factory is structured, or make the customer do it. The table-stakes copilot is commoditized.

The empty chair = the contextualization layer (the BEFORE)

Our moat is the bundle no one has whole: cited + self-scored + read-only/UNS-safe + works-on-your-mess. The window is ~12 months before a MaintainX-class incumbent adds citations + scoring. Our sharpest, least-contested knife: per-answer, customer-visible groundedness scoring - only Cognite tells a benchmark story, and theirs is a once-a-year PDF. Ours is live, per answer.

3. Where we really are - the honest read

- Zero paying customers. The 90-day MVP's 'first 3 logos' window has passed with none.

- The M
CI-en
and th

The wedge is real but not yet earned. The window is ~12 months. Provelt 2027 is the forcing function.

4. The way forward - one storyline

Make the loop 'contextualize a factory we didn't seed -> cited diagnosis' actually work, prove it on someone else's mess, and let Provelt! 2027 (Feb 9-12, Dallas) be the deadline that forces it.

- Close the contextualization loop on FOREIGN data - generalize the parser/tag-mapper past the bench; enforce citations; connect read-only to a foreign UNS.

- Prove
'stran

5. Provelt! 2027 - the perfect conference

Provelt! (Feb 9-12, 2027, Hilton Anatole, Dallas; Walker Reynolds / 4.0 Solutions) is NOT a scored benchmark. It is a live showcase: every vendor connects to the SAME intentionally-messy shared UNS factory and answers four questions on stage - what problem, how, how long, how much - judged by the manufacturers in the room. 2026's energy: UNS + knowledge graphs + agents building live with Claude Code. No winner.

The 20-minute arc

- Frame: 'Everyone here shows an agent that works AFTER the factory is contextualized. We'll do the part everyone assumes away - live, on a factory we've never seen.'

- Conte
knowl
trustw

Closing line: 'Everyone here showed you an agent that works after your factory is contextualized. We just contextualized a messy factory we'd never seen - live, in twenty minutes - and answered a real fault with a cited source. That's the part everyone else assumes is already done.'

The prerequisite checklist for the stage = the beta gate. If we can contextualize Provelt's shared messy factory live, the beta gate is closed by definition. Same finish line. (Full runbook: docs/plans/2026-06-22-provelt-2027-demo-runbook.md.)

6. Decisions - locked and open

Locked

- The context layer is the wedge (FactoryLM), the agent is the proof (MIRA) - not a copilot, not a platform we rebuild. Lead with context.

- Produ

Open - your call

- Pricing: pick ONE architecture (the docs carry three) before beta outreach.

- The f
share