

FactoryLM Contextualized-Diagnosis Proof Packet

Filler 01 - Underfill from Low Bowl Pressure

VERDICT: PASS | scenario_id: filler_underfill_low_bowl_pressure | tenant (SimLab): 00000000-0000-0000-0000-000000515ab1
| env: staging (factorylm/stg) | generated 2026-06-24

1. Executive Summary

A fault was injected into the SimLab juice-bottling line on **filler01**. Live signals were landed through the deployed relay ingest path into `live_signal_cache`, and the real MIRA Supervisor was asked the operator's question. MIRA returned a **grounded, cited diagnosis** (root cause + live evidence + manual citation + corrective action), not a clarifying question. All evidence below is independently verifiable against the staging database.

2. Technical Diagnosis

Field	Value
Fault	Filler 01 - Underfill from Low Bowl Pressure
Affected asset	filler01
UNS path	enterprise.florida_natural_demo.plant1.juice_bottling.line01.filler01
UNS source	direct_connection (certified — chat-gate skipped)
Expected root cause (ground truth)	Low filler bowl pressure causing underfill
Expected actions (ground truth)	Check compressed-air header pressure at AirSystem01; Inspect and adjust filler bowl pressure regulator; Verify fill-valve air supply manifold; Inspect nozzle for clogging
MIRA confidence mode	DIRECT_ANSWER (commit, no Socratic question)

3. Live Tag Evidence (observed vs normal)

Tag	Observed (faulted)	Normal baseline
filler_bowl_pressure	5.2998	12.0
fill_level_oz	12.7547	16.0
underfill_reject_count	7	0

Ingest evidence: relay POST /api/v1/tags/ingest -> **89 accepted, 0 rejected**; tag_events rows=89 (distinct 89, all_simulated=True); live_signal_cache rows=89.

4. Retrieved Documents & Citations

#	Retrieved chunk (recall_knowledge, SimLab tenant)
1	simlab://florida_natural_demo/line01/rinser01/plc_tag_description_sheet.md # Rinser 01 - PLC Tag Description Sheet ## Process Tags Tag Name Full UNS Path Type Unit Normal Range Description ----- ----- ----- ----- ----- ----- -----
2	simlab://florida_natural_demo/line01/casepacker01/troubleshooting.md # Case Packer 01 - Troubleshooting Guide ## Jam Detected - Line Backs Up Upstream **Condition:** `jam_detected` = TRUE; `run_state` = HELD or ABORTED; upstre
3	simlab://florida_natural_demo/line01/filler01/troubleshooting.md # Filler 01 - Troubleshooting Guide ## Underfill / Low Fill Weights **Condition:** `fill_level_oz` consistently below `fill_level_target_oz`; `underfill_reje
4	simlab://florida_natural_demo/line01/filler01/fault_code_table.md # Filler 01 - Fault Code Table Code Label Severity Description Likely Cause Recommended Action ---- ---- ----- ----- ----- ----- -----
5	simlab://florida_natural_demo/line01/airsystem01/operator_quick_guide.md # Air System 01 - Operator Quick Guide ## System Overview Air System 01 is the plant compressed-air utility for Line 01. It supplies the entire juice bottling

Citations emitted in MIRA's answer: intro, Process Tags, Process Tags, intro, Process Tags, intro, intro, intro, Process Tags

Expected citations (ground truth): troubleshooting.md; fault_code_table.md; operator_quick_guide.md

5. MIRA Answer (verbatim)

Q: Why is Line 1 making bad bottles?

A: The filler01 is currently in an Idle state and has a fault code of F010, which is not defined in the provided reference documents [Source: intro]. The `fill\_level\_oz` is 12.7547, which is lower than the target fill level of 16.0 oz [Source: Process Tags]. The `filler\_bowl\_pressure` is 5.2998 PSI, which is below the normal range of 10-18 PSI [Source: Process Tags]. The `underfill\_reject\_count` is 7, indicating a systematic underfill issue [Source: intro]. Check the `tank\_level\_percent` to ensure it is not low, as this could be causing the underfill issue [Source: Process Tags].
2. Inspect the pressure regulator to ensure it is not blocked or failed [Source: intro].
3. Verify the AirSystem01 header pressure to ensure it is not low, as this could be affecting the pneumatic actuators [Source: intro].
4. Check the nozzles for any blockages or issues [Source: intro].
5. Adjust the `fill\_level\_target\_oz` and `vfd\_speed\_hz` to optimize the fill level and production rate [Source: Process Tags].

6. Database Evidence — live_signal_cache (faulted asset)

uns_path	value	sim	src
filler01.faults.fault_code	F010	True	simulator
filler01.faults.nozzle_fault_count	0	True	simulator
filler01.motor.motor_current_amps	8.4692	True	simulator
filler01.motor.vfd_speed_hz	46.1013	True	simulator
filler01.process.bottles_per_minute	0	True	simulator
filler01.process.fill_level_oz	12.7547	True	simulator
filler01.process.fill_level_target_oz	16	True	simulator
filler01.process.fill_level_variance	0.6	True	simulator
filler01.process.filler_bowl_pressure	5.2998	True	simulator
filler01.process.product_temperature	38.4712	True	simulator
filler01.process.tank_level_percent	74.1301	True	simulator
filler01.quality.overfill_reject_count	0	True	simulator
filler01.quality.underfill_reject_count	7	True	simulator
filler01.status.run_state	Idle	True	simulator

Verify yourself (staging, read-only): `SELECT uns_path, last_value_numeric, last_value_text FROM live_signal_cache WHERE tenant_id='00000000-0000-0000-0000-000000515ab1' AND uns_path::text LIKE '%filler01%';`

7. Langfuse Observability

Field	Value
Trace name	supervisor.process
Trace ID (this run)	(not emitted — see note)
Where to find it	Langfuse instance at \$LANGFUSE_HOST -> Traces -> filter name='supervisor.process', user_id=chat_id
Spans to inspect	the recall/retrieval span (input: query; output: KB chunks) and the LLM generation span (input: system prompt + chunks +
Inputs/outputs to verify	generation INPUT contains the cited chunk text; generation OUTPUT cites that chunk -> grounded
Grounded vs hallucinated	GROUNDED if the answer's facts (e.g. the value + normal range) and the [Source:] citation both appear in the retrieved cl

Honest note: this proof run executed the engine directly under Python 3.14, where Langfuse cannot emit (langfuse v2 = the telemetry-compatible client = incompatible with Py3.14; langfuse v4 dropped the .trace() API). So NO live Langfuse trace exists for THIS run. The DEPLOYED engine (Py 3.12 + langfuse 2.50 + \$LANGFUSE_HOST) emits the 'supervisor.process' trace per turn. Grounding here is instead proven independently by sections 3-6 (the DB-backed retrieved chunks vs. the answer's citations) — which does not require trusting a trace.

8. Pass / Review Verdict (graded by simlab.diagnostic.grade)

Rubric dimension	Result
------------------	--------

root_cause_hit (expected phrase matched)	PASS
asset_hit (expected asset named)	PASS
evidence_recall (>=0.5 to pass)	0.75
citations_hit (by filename)	(none - answer cites doc SECTIONS, not filenames)
actions_hit	Inspect and adjust filler bowl pressure regulator
rubric_passed	PASS

Verdict rule: rubric_passed AND ingest==200 AND no_trailing_question. **REVIEW** is reported honestly, not hidden — a partial match (e.g. the answer is reasonable but does not contain the exact ground-truth root-cause phrase, or cites a doc SECTION rather than its filename) shows here as a failed dimension.

Overall: PASS

9. Provenance & Reproducibility

Field	Value
git SHA / branch	7c1433ca3351 (docs/proof-packets) [DIRTY working tree]
Config (env)	{"MIRA_TENANT_ID": "00000000-0000-0000-0000-000000515ab1", "MIRA_SHARED_TENANT_ID": "00000000-0000-0000-0000-000000515ab1", "MIRA_DIRECT_ANSWER_MODE": "1", "MIRA_RETRIEVAL_HYBRID_ENABLED": null}
Corpus	242 chunks, doc-set md5 977b0f975778
Retrieval mode (actual)	bm25-only (no vector/rerank - embedder/reranker down)
Embedder / reranker / langfuse	ollama_reachable=True, nvidia_configured=True, langfuse=False (py 3.14.2)
Reproduce	doppler run --project factorylm --config stg -- python tools/proof/run_proof.py --clean (seeds: tools/seeds/seed-simlab-docs.py)