HubV3 User Test Sheet prod v3.29.0 · shipped 2026-06-20

Contextualization intake — **upload evidence → Hub stages it → a human reviews → approve → it goes live.** Three ingest surfaces (offline app, Hub, Telegram) + automated suite. Companion: `hubv3-rollback.md`, `garage-conveyor-demo.md`.

Tester: ________________ **Date:** ________________ **Build:** ________________

> **What you are proving:** nothing a user uploads reaches MIRA's answers until a human approves it. Test A = evidence in. Test B = Hub owns the truth (review→approve→publish). Test C = the guarantees that make it safe. Test D = fast automated proof. **At ship, only A & D were run live — B & C need an authed Hub session and are unconfirmed on prod. You are the one confirming them.**

TEST A — Offline Contextualizer

local desktop app · ✅ verified at ship

> Start it (the venv carries the PDF/OCR deps — launching any other way makes A4 fail):

```
cd /c/Users/hharp/Documents/MIRA-pr2068/mira-contextualizer && .venv/Scripts/python.exe -m mira_contextualizer
```

> An Edge app-window opens on localhost.

#	DO	EXPECT	PASS / FAIL
A1	File → New profile → "Garage Demo / Micro820 Conveyor"	empty profile + identity editor	☐ ☐
A2	Import Files → `MIRA\plc\Micro820_v4.1.9_Program.st`	status **done**, ~62 signals in Extracted Signals	☐ ☐
A3	Import Files → `MIRA\plc\MbSrvConf_v4.xml`	Modbus tags added, no dup of A2	☐ ☐
A4	Import Files → `Downloads\gs10usermanual.pdf`	**done** (not "error") — wait ~1 min, 453 pages	☐ ☐
A5	Click UNS Map / Fault Catalog (ISO 14224) / Parameters (UCUM) / Scorecard	populated, not empty	☐ ☐
A6	**Export Bundle** (full)	`machine_context_bundle.zip` downloads	☐ ☐
A7	**Export Bundle** → sanitized mode	smaller zip; **no** `documents/*.json` raw payloads, but `evidence.json` + `uns.json` present	☐ ☐

PASS = A4 done (not error) + A6/A7 both produce zips + A7 has no raw docs.

Common failures: A4 "error: No module named 'pdfminer'" → app launched without the `[docs]` extra; use the venv path above. CCW `.ccwsln` imports as "other" with 0 signals → that's a solution-pointer file; import the `.st` / `MbSrvConf.xml` instead.

Notes:

TEST B — Hub import → review → approve

app.factorylm.com/hub · ⚠ unconfirmed on prod

> The core "Hub owns truth" flow. Be logged in at `https://app.factorylm.com/hub` .

#	DO	EXPECT	PASS / FAIL
B1	Find **Contextualization** / **Import Review** in the sidebar	nav link present (P7)	☐ ☐
B2	Import the `machine_context_bundle.zip` from A6	creates an **import batch**, status `proposed`	☐ ☐
B3	Open the **Review Queue**	staged signals / UNS / faults / params, all **pending**; shared labels (Sources, Evidence, Extracted Signals, Fault Catalog, Parameters, UNS Map, Scorecard)	☐ ☐
B4	Re-import the **same** zip	**no duplicate** sources (sha256 dedup)	☐ ☐
B5	**Approve** the batch	staged items publish; `kg_entities proposed → verified`, available to MIRA	☐ ☐
B6	Re-import after approve	does **not overwrite** approved data — skip reason shown, not silent	☐ ☐

PASS = B2 creates a batch, B4 no dup, B5 publishes, B6 no overwrite. This is the whole product thesis.

Notes:

TEST C — The guarantees ("train before deploy")

⚠ unconfirmed on prod

#	CHECK	EXPECT	PASS / FAIL
C1	**Nothing auto-promotes:** between B2 and B5, read the staged items	everything reads `proposed` — MIRA cannot answer from it yet	☐ ☐
C2	**Asset matching:** import against existing conveyor asset / brand-new machine / ambiguous	existing → stages **under it** (strong match); new → **draft asset proposal**; ambiguous → **needs confirmation**	☐ ☐
C3	**Telegram (if wired):** send a nameplate photo + caption to the bot	lands in the **same Review Queue** as `proposed` (ingest_route=telegram); does **not** bypass to the KB	☐ ☐

PASS = staged data stays `proposed` until B5; matching routes correctly; Telegram cannot skip review.

Notes:

TEST D — Automated proof

fast, no UI · ✅ verified at ship

The §6 acceptance matrix — re-run anytime:

```
cd /c/Users/hharp/Documents/MIRA-pr2068/mira-hub && bun run test src/lib/contextualization/
```

#	SUITE	EXPECT GREEN	PASS / FAIL
D1	asset-matcher	11 pass	☐ ☐
D2	acceptance-matrix	11 pass	☐ ☐
D3	approval	12 pass	☐ ☐
D4	intake-contract	8 pass	☐ ☐
D5	bundle-import / promote / routes	5 / 6 / all pass	☐ ☐

PASS = every suite green. Scripted end-to-end demo: `docs/runbooks/garage-conveyor-demo.md`.

Notes:

If it breaks → rollback

SYMPTOM	LAYER	ACTION
`/hub/api/contextualization` **500s** (not 401)	DB / migration	**ROLLBACK B** — `056` down-SQL or Neon `br-square-cake-ah54ijqu`
Hub won't load / bad bundle after deploy	deploy	**ROLLBACK A** — redeploy `checkpoint/pre-hubv3-2026-06-20`
Code must leave main	git	**ROLLBACK C** — revert the merge

Healthy prod check: `/hub/api/contextualization` returns **401** (alive + auth-gated, tables present). Full procedure: `docs/runbooks/hubv3-rollback.md`.