

70-Day Sprint to XPRIZE Submission

FactoryLM / MIRA — Build with Gemini XPRIZE

Prepared June 08, 2026 — Confidential

***North Star:** By August 17, a stranger can sign up for a FactoryLM Maintenance AI Readiness Assessment, pay \$500–1,500, receive an AI-generated report grounded in their real equipment data, and MIRA can answer their maintenance questions with cited evidence. Revenue is real, customers are real, AI is in production.*

Three Scoring Pillars (Equal Weight)

Pillar	Strategy
Business Viability (33%)	Real revenue during the 90-day window. Arms-length customers. Sustainable model. Target: 2–4 paid DT Assessments (\$500–1,500 each) + 1 MIRA Pilot (\$2–5K/mo) = \$2,500–11,000 total revenue.
AI-Native Operations (33%)	AI is live in production, executing key decisions continuously. Already strong: 14 scheduled AI agents, inference cascade, orchestrator pulse, automated lead pipeline, eval framework. Polish the evidence.
Category Impact (33%)	How significantly the software moves the needle in Small Business Services. The pitch: 80% of small manufacturers can't get maintenance intelligence support. MIRA gives a \$50M plant the same diagnostic capability as a Fortune 500 reliability team.

Week-by-Week Sprint

Week 1 (Jun 8–14): Register + Revenue Foundation

- Register at geminixprize.com as FactoryLM (Organization)
- Claim \$100 AI Ultra + \$300 Cloud credits
- Share MIRA repo with testing@devpost.com + judging@hacker.fund
- Activate Stripe LIVE mode (currently test mode) — connect to Wells Fargo or backup card
- Publish DT Assessment landing page on factorylm.com with Stripe checkout (\$500 tier)
- Wire Gemini as a visible, documented call in the production inference path

- Fix the 3 beta blockers from orchestrator pulse (rate limit, tenant scope, safety eval)

Week 2 (Jun 15–21): First Customer Sprint

- LinkedIn reactivation: 3 posts targeting the 4K hydraulic audience — "what I wish I had" stories
- Direct outreach to 10 HubSpot leads (Tampa + Florida pipeline) offering free 30-min consult + paid assessment
- Reach out to Shingle & Gibb (expo follow-up) and Aercon Florida (discovered lead)
- SimLab juice bottling line demo working end-to-end (in progress — should complete this week)
- Record rough-cut demo video showing MIRA answering "Why is Line 1 making bad bottles?"

Week 3–4 (Jun 22 – Jul 5): Close First Sale + Polish Product

- GOAL: Close first paid DT Assessment (\$500–1,500)
- Deliver the assessment using MIRA + manual analysis — document the AI-assisted workflow
- Get written testimonial from first customer
- Apply the quickstart rate-limit patch + tenant-scope fix to production Hub
- Run the SimLab underfill scenario through the full MIRA diagnostic chain — record screenshots
- Train-before-deploy: demonstrate the approval flow with real SimLab data

Week 5–6 (Jul 6–19): Scale Revenue + Build Evidence

- GOAL: 2–3 more paid assessments or 1 MIRA Pilot subscription
- LinkedIn content cadence: 2 posts/week showing real MIRA capabilities
- Collect Stripe dashboard exports, start the P&L; document
- Document all agent execution logs — orchestrator pulse history, lead discovery runs, eval regressions
- Screenshot the Hub dashboards: knowledge graph, command center, asset agents
- Begin writing the 500–1,000 word narrative (human vs AI, jobs created, building story)

Week 6–7 (Jul 20 – Aug 2): Record Demo Video

- Script the 3-minute video: (1) problem statement, (2) live MIRA demo, (3) revenue/customer evidence
- Film MIRA answering a real maintenance question with cited evidence from a customer's equipment
- Show the AI-native operations: orchestrator pulse, lead pipeline, automated eval
- Show the SimLab bottling line as a Provelt demo

- Upload to YouTube (public)

Week 8–9 (Aug 3–14): Final Submission Assembly

- Finalize revenue evidence (Stripe export, bank statement, P&L;)
- Finalize customer evidence (names, emails, testimonials with permission)
- Finalize the written narrative
- Verify GitHub repo access for judges
- Compile product evidence: agent logs, API usage, dashboard screenshots
- Document expenses: DigitalOcean, NeonDB, Doppler, domain, Gemini API costs
- Internal review: does submission score well on all 3 pillars?

Week 10 (Aug 14–17): Submit

- Final submission on xprize.devpost.com by August 17, 1:00 PM PT
- Double-check all required fields: video link, repo link, narrative, revenue, expenses, customer evidence
- Verify video is public on YouTube and under 3 minutes
- Submit and celebrate — regardless of outcome, you now have a revenue-generating product

Revenue Model for the 90-Day Window

The competition grades actual revenue, not projections. Here are three scenarios:

Scenario	Revenue Sources	Total	Competition Impact
Minimum Viable	1 DT Assessment (\$500) + 0 pilots	\$500	Gets on the board. Weak but real.
Target	2 Assessments (\$500 ea) + 1 Pilot (\$2K/mo x 2mo)	\$5,000	Competitive for runner-up prize.
Stretch	3 Assessments (\$1,000 avg) + 1 Pilot (\$3K/mo x 2mo)	\$9,000	Top 5 contender if demo is strong.

Fastest Path to First Dollar

The DT Assessment is the fastest monetizable unit. It requires no software deployment at the customer site — just their equipment list, photos, and maintenance pain points.

Timeline	Action
Day 1–3	Publish assessment landing page with Stripe checkout on factorylm.com
Day 4–7	Direct outreach to 10 warm leads (HubSpot pipeline + expo contacts)
Day 7–10	Free 30-min consult calls — demonstrate MIRA on their equipment type
Day 10–14	Close first paid assessment (\$500)
Day 14–21	Deliver assessment report (AI-generated + manual QA)
Day 21+	Upsell to MIRA Pilot (\$2–5K/mo) or second assessment

Gemini Integration Strategy

The competition requires at least one Gemini API call in the deployed application. MIRA already has Gemini in the cascade. To make this maximally visible to judges:

- **Document the cascade in the repo README:** Groq → Cerebras → Gemini with explicit mention that Gemini handles diagnostic reasoning when upstream providers are unavailable
- **Add Gemini-specific logging:** tag every Gemini API response in the agent execution logs so judges can see it in the product evidence
- **Use the \$100 AI Ultra credits** for the SimLab diagnostic demo — show Gemini reasoning over bottling line fault scenarios

- **Optional:** Use Google Cloud Run or Cloud Functions for one SimLab route to satisfy the "at least one product from Google Cloud" requirement (currently using DigitalOcean VPS)

Risk Register

Risk	Level	Mitigation
No revenue by Aug 17	HIGH	Submit anyway — score will be weak on Business Viability but AI-Native Ops and Category Impact may still earn a runner-up. Meanwhile: the assessment pipeline and LinkedIn reactivation start Week 1.
Wells Fargo card keeps bouncing	HIGH	Get a backup card (personal CC or prepaid) before Stripe, DigitalOcean, Neon, and GCP all suspend simultaneously.
Gemini rate limits / 403s	MEDIUM	The cascade falls through to other providers. Document this resilience in the submission — it's a feature, not a bug.
SimLab not ready in time	LOW	The garage conveyor + real PLC bench already work as a demo. SimLab juice bottling is bonus, not required.
Competition too crowded	LOW	2,700+ teams but almost none in industrial maintenance. Small Business Services for manufacturing is a blue ocean in this hackathon.

Why MIRA Stands Out in This Competition

Most XPRIZE entries will be AI wrappers — chatbots, content generators, SaaS dashboards with GPT bolted on. MIRA is fundamentally different:

- **Domain depth:** 30 years of maintenance experience encoded into the product. Judges can tell the difference between a weekend project and a domain expert's life work.
 - **AI is the business, not a feature:** 14 autonomous agents run FactoryLM's operations. The company literally cannot function without AI. That's the AI-Native Operations score.
 - **Real industrial hardware:** PLC bench, Ignition gateway, Modbus comms, live conveyor. This isn't a mockup — it's a working factory cell.
 - **UNS architecture:** ISA-95 compliant namespace that Walker Reynolds himself would recognize. No other hackathon entry will have this.
 - **Train-before-deploy:** Human-in-the-loop approval that judges will love — responsible AI in an industry where getting it wrong can hurt people.
-

Prepared by the FactoryLM AI operations team for Mike Harper. 70 days to submission. The clock is running.